Reproducibility of data analysis

Stephanie Palmer & Stefano Allesina & Graham Smith

September 12-14, 2017

Contents

Goals

This tutorial will cover methods for making your code reproducible, both by you and by others. In particular, we will discuss methods to ensure that you can: reproduce your calculations precisely, make your code readable, track changes to your code, and make your code freely accessible to the wider world via github. Along the way, you will get an introduction to stochastic processes and how they are used to model biological variability. By the end of this tutorial, you should know why it is important to save your seeds and merge your branches.

Installation notes

For this tutorial, relatively up-to-date versions of `R` and `RStudio` are needed. To install `R`, follow instructions at `cran.rstudio.com`. Then install `Rstudio` following the instructions at `goo.gl/a42jYE`. Download the `ggplot2`, `cowplot`, `stats`, and `RMKdiscrete` packages. You should also install the UNIX utility `git`.

For the bonus data challenge, you will also need to install the Unix utility `libcurl`. If you haven't done so already, you can find installation instruction links here for OSX and Windows.

Commenting your code

When you begin programming, you have written just a few programs and functions and it might be possible to hold all of your naming conventions, subroutine purposes, and little tricks in your mind at once. However, wait just a year or even just a few dozen code snippets later and you'll have a hard time deciphering, perhaps even running *your own* code if you don't take the time to comment your files carefully. Best practices in modern computational science also dictate that you share your code whenever you publish a result. This means that other people need to be able to read your code; it should run on other computers with other folks sending the input data and specifying or *changing* parameters. This might seem daunting, but a few simple habits will aid in keeping all of your code parsable by other humans, including your own future self.

Include a comment block at the top of your program, function, script, etc. that describes who wrote the code, when it was last updated, and its input/output functionality, and generally, the code's purpose. For

advanced users, there are packages, e.g. `roxygen2`, that will automatically convert text in this block into `.Rd` documentation in the `man` folder, that can then be accessed by typing `?` or `help()`.

Code comments should describe what you *intend* the code to do, not what it does in detail. Comments that do the latter are simply restating in words, and probably in many more characters and more clumsily, what the actual code is saying. This is redundant; don't do it.

Avoid the need to write lots of explanatory commments by writing Really Obvious Code (ROC). Tips for writing ROC: * use function and variable names that are self-explanatory. For example `random_locations_of_N_spiders_in_a_box.R` is a much better function name than `eek.R`, and `eek2017.R` is still better than `eek.R`. Old coders had to limit variable name and code lengths because of the limitations of memory size. You're no longer constrained in this way!

DIY lottery

To practice commenting code, let's write a short program to generate a draw from a lottery. The lottery is a somewhat peculiar one: The winning lottery number is drawn by flipping a coin 30 times and reporting a "1" for heads and a "0" for tails. The same bent coin is always used, and it has a probability of 0.1 of turning up heads and a probability of 0.9 of turning up tails. Each ticket for the lottery is a buyer-determined string of 30 1's and 0's, and costs \$1. The prize money for the lottery is \$10,000,000.

Let's make sure to load the `stats` package.

```r
library(stats)
```

Next, let's explore how we simulate flipping this bent coin once. We're going to use one of R's random number generator functions `runif`, which will give you uniformly distributed numbers on the interval [0.1]. Here's one way to do it:

```r
# define the probability of heads
p_heads <- 0.1
# draw a random number from the interval [0,1]
random_val <- runif(1)
# threshold the random number based on p_heads to decide the sign of the flip
coinflip <- if (random_val < p_heads) 1 else 0
# 1 is for heads, 0 is for tails
```

> **Exercise**
>
> Write a function in R that will simulate one random draw for the winning lottery ticket. Let the inputs to the function be the number of digits in the lottery ticket (let's use 10 to start) and the probability of heads (keep this at $p = 0.1$ to start). The function should output a sequence of 1's and 0's. Use all of the best practices for commenting your code.

Each flip of a coin like this with probability, p, of heads is an example of a Bernoulli trial, the general term for an experiment with only two output states, success or failure. The number of heads in the sequence of independent coin flips generated by our lottery will follow a binomial distribution

$$P_n(k) = \binom{n}{k}p^k(1-p)^{n-k},$$

where p is the probability of heads (1's), n is the length of our lottery ticket, and k is the number of heads in the ticket. The prefactor $\binom{n}{k} = \frac{n!}{k!(n-k)!}$ is called the binomial coefficient and describes the number of unique ways of placing k identical objects in n bins. Capital "P" represents the probability distribution of k heads.

You can use the function `rbinom` to generate draws from a Bernoulli distribution. In our lottery, this would amount to flipping the coin n_{flips} times.

```r
n_flips <- 3
rbinom(n_flips,1,p_heads)
```

```
# [1] 0 0 0
```

Exercise

If you could only buy one ticket for this lottery, which one would you buy? Modify your code (perhaps use `rbinom`) and run it several times (maybe 100 or even 1000) to sample from lots of possible lottery outcomes. Did you pick a good number? How many times did your number come up?

Save your seeds

Compare a draw from your lottery with your neighbor. Do you draw the same sequence of random lottery tickets? Why not? If you wanted to reproduce the *exact* same output from your lottery each time you decide to reset it, you'll need to know a little more about how R's random number generator (RNG) works. Try typing:

```
? RNG
```

You will notice that the function `RNGkind` is the interface for querying the current state of the RNG. Let's find out what the current settings are:

```
RNGkind()
```

```
# [1] "Mersenne-Twister" "Inversion"
```

The first part is the RNG algorithm, the second specifies the algorithm for transforming uniformly distributed random numbers into random samples from the normal, or Gaussian, distribution. The twister algorithm based on Mersenne prime numbers, $M_n = 2^n - 1$, where n is also prime, is a state-of-the-art pseudo-random number generation scheme, developed by Matsumoto and Nishimura in 1997. NB: RNG's should technically be called *pseudo*-random number generators or PRNG's, in part because they all have some period after which they will produce exactly the same sequence. The trick is to find an algorithm with a period so long you'll never notice the "P" in the "PRNG". The Mersenne Twister algorithm has a period of $2^{19937} - 1$ and passes many statistical tests for randomness.

The seed to an RNG is usually a large integer that provides an initialization the RNG algorithm. Scrolling down to "Note" you'll learn that the seed to R's RNG is set by the current time and sometimes the process ID. That means that your simulation results will depend on when you start your R session, run your code, and even some local information in your processing environment. Compare the output of `runif` with your neighbor.

```
runif(5)
```

Starting an RNG with the same seed will produce exactly the same sequence of random numbers; an RNG spits out random numbers, but not noisy ones. To reproduce your simulation results precisely when you use an RNG, you'll want control of that seed. R uses `set.seed` which takes a small integer as input and generates, deterministically, all the random seeds necessary for your RNG algorithm.

```
set.seed(19937)
runif(5)
runif(5)
set.seed(19937)
runif(5)
```

Exercise

Test whether or not you and your neighbor get precisely the same sequence of numbers when you use the same seed. Will this work if you use different RNG algorithms? Try changing your RNG algorirhm using `RNGkind` and compare your results with your neighbor, when you use the same

seed. If you wanted to be able to instruct someone to reproduce your exact simulation results using `R`'s RNG, what would you need to tell them?

You can use `RNGkind` to set or query both the RNG and normal algorithms. You can save this information along with the current value of `seed` using something like:

```
seed <- 19937
set.seed(seed)
seed_used <- seed
RNGkind_used <- RNGkind()
save("seed_used","RNGkind_used",file=RNGinfo_for_mycode)
```

When you are ready to share your code with others, you should also save all the version information for your current `R` package and libraries to this same file.

Noise in biological systems

Many of the variables that we observe in biological recordings fluctuate, sometimes because we cannot control all the states of the external and internal experimental system, other times because thermal noise makes the state of the biological system we interrogate inherently variable. Examples of fluctuating quantities in biological systems include: the number of a certain type of molecule in a cell; the number of open channels in a cell; the number of electrical action potentials or "spikes" emitted by a neuron in response to a stimulus; the number of individuals in a population at a particular moment in time; the number of bacterial colonies on a plate. These are all quantities that we can make precise claims about, on average, but cannot specify with certainty for any particular experimental observation.

It is useful to model not only a mean value for a fluctuating variable, but the full shape of its distribution of values. For example, if we observe the firing of neurons in the brain to repeats of the same external stimulus, the precise times of spikes will vary between repeats. By fitting the statistics of this noise to models, we deepen our mechanistic understanding of the neural response. We can test whether or not the "noise" we observe is consistent with a truly random source of output variation, or if it has some structure that tells us about interactions between the neurons and their environment.

Often, noise in biological systems is modeled by what's called a Poisson process, whose values follow a Poisson distribution. The program you wrote to sample the bent coin lottery generated tickets whose statistics follow the binomial distribution. The binomial distribution approches the familiar Poisson distribution, in the limit of a large number of trials, n, or a small probability of the event, p, per trial:

$$P_n(k)_{n \to \infty} = \frac{\lambda^k}{k!} e^{-\lambda}$$

where λ is the average rate of occurrence of our event in n trials. We have just written down the Poisson distribution. You will see this used as a model for biological variability again and again, either explicitly or implicitly. It is important to think about whether or not it is a good model for the system under study each time you come across it or are deciding to use it for your own research. Notes on deriving the relationships between some common distributions are provided in the readings folder.

Arthropod dispersion

In 1941 and 1942, zoologist LaMont Cole, then working at the University of Chicago, set out to survey species diversity and distribution in the woods and pastures of Kendall County, Illinois. He was particularly interested in which species co-occurred in woods versus grazing land and how their numbers varied with changes in humidity and temperature throughout the year. He laid thick oak boards in a variety of locations and counted the number of "cryptozoic" (animals found under stones, rotten logs, tree bark, etc.) individuals found under the boards several times a week, over the course of a year. For his spatial distribution studies, he aimed to determine whether arthropods distributed themselves randomly or if they had a more complex interaction pattern with each other or with their environment.

Spiders, spiders everywhere

Load the arthropod data located in the data directory for the "Reproducibility" tutorial. If your current directory is `code`:

```
x <- read.delim("../data/cole_arthropod_data_1946.csv", sep=',')
```

Have a look at the data:

```
head(x)
```

```
#   k_number_of_arthropods C_count_of_boards_with_k_spiders
# 1                      0                             159
# 2                      1                              64
# 3                      2                              13
# 4                      3                               4
# 5                      4                               0
# 6                      5                               0
#   C_count_of_boards_with_k_sowbugs
# 1                               28
# 2                               28
# 3                               14
# 4                               11
# 5                                8
# 6                               11
```

We have three columns of data that correspond to k arthropods found on a board, and the count, C, of the number of boards found at the site with k arthropods. For example,

```
x$k_number_of_arthropods[2]
x$C_count_of_boards_with_k_spiders[2]
```

```
# [1] 1
# [1] 64
```

tells you that 64 boards were found with 1 spider under them. Let's have a look at the total number of boards:

```
N_boards <- sum(x$C_count_of_boards_with_k_spiders)
```

Next, let's compute the total number of spiders and the mean number of spiders per board:

```
N_spiders <- sum(x$C_count_of_boards_with_k_spiders*x$k_number_of_arthropods)
spider_rate <- N_spiders/N_boards
print(spider_rate)
```

```
# [1] 0.425
```

Finally, let's plot the distrubtion of observed spiders per board. We'll want to make sure we express this as a probabilty of finding k spiders on a board, so we'll have to divide by the total number of boards.

```
x$P_k_spiders <- x$C_count_of_boards_with_k_spiders/N_boards
```

Load the `ggplot2` and `cowplot` packages if you haven't already, and make a plot of $P(k)_{\text{spiders}}$.

```
library(ggplot2)
library(cowplot)
```

```
g_spiders <- ggplot(x, aes(x=k_number_of_arthropods, y=P_k_spiders)) + geom_point(size = 3) + xlab("k sp
g_spiders
```

Now, compute the Poisson distribution of expected spiders per board if the spiders have a rate `spider_rate` and are distributed randomly across the boards.

```
lambda = spider_rate
k <- 0:17
Poisson_k_spiders <- data.frame(k=k, probability=lambda^k*exp(-lambda)/factorial(k))
```

You can also use the built-in function `dpois` from the `stats` package to generate the probabilities in `Poisson_k_spiders`.

```
Poisson_k_spiders <- data.frame(k=k, probability=dpois(k,mean_spiders))
```

Add a line for the Poisson model to your plot of the observed spider probability.

```
g_spiders <- g_spiders + geom_line(data=Poisson_k_spiders, aes(x=k, y=probability), linetype='dashed',
```

Exercise

How well does the Poisson model account for the observed distribution of spiders? How might you quantify this?

If the spider rate were closer to 7 spiders per board, what would you expect for the distribution of k spiders per board? Let's plot this distribution.

```
Poisson_k_spiders_mean_7 <- data.frame(k=k, probability=dpois(k,7))
ggplot(x, aes(x=k_number_of_arthropods, y=P_k_spiders)) + geom_line(data=Poisson_k_spiders_mean_7, aes(
```

Does this distribution look familiar?

Simulating dispersion

We know the equation for the Poisson distribution, so it was an easy task to estimate the rate of finding spiders on a board and compare our observed results to the known distribution. What if we had an underlying microscopic model that didn't have such a derivable formula? Let's explore how to simulate a random process by checking that we do in fact generate samples from a Poisson distribution when spiders are located at random positions in an area.

Exercise

Write code that randomly places 42 spiders in a 10×10 area. Each 1×1 box is a "board" and you have 100 boards in your simulation. Count the number of spiders per board for one run of your simulation. How does your P_k distribution compare to the Poisson distirbution with rate $p = 0.42$? How could you improve your simulation? Think about running many experiments or sampling a larger area. Are these equivalent? Use what you have learned about commenting your code as you write. Finally, swap code with your neighbor and see if they can run and plot your results.

Sowbugs are social

Cole also counted the numbers of many other species he found under the boards. He noticed that sowbugs showed an interesting deviation from Poisson behavior. Note: sowbugs or pillbugs are land-dwelling crustaceans, more closely related to shrimp than to insects.

Exercise

Using what you have learned in the spiders section, compute the average rate of sowbugs per board, plot the observed probability of finding k sowbugs per board, and plot the Poisson distribution with the observed sowbug rate. How well does the Poisson distribution match up with the observed sowbug distribution?

The lack of fit of the Poisson model shows that sowbugs may be more social than spiders; they are frequently found in large clusters - quantitatively more so than would be expected by chance, given their density. This is an exmaple of how comparing biological variability to what is expected given certain simple assumptions (such as "arthropods don't notice each other") can lead to biological insight.

Cole hypothesized that sowbugs are attracted to one another, and prefer to aggregate into groups of 3 or more individuals. He further hypothesized that these aggregates disperse themselves randomly. The formation of the aggregate is a random process, as is the dispersion of the aggregates. This is an example of a doubly-stochastic process.

In the late 1970's, Janardan and colleagues showed that the count distribution of sowbugs and other organisms with social interactions can be well-described by a modified Poisson process, a process resulting in a Lagrangian Poisson Distribution or LGP, that accounts for the attraction or repulsion of individuals to each other. The LGP has a mean rate parameter, λ_1, just like in the Poisson distribution. The LGP has a second parameter, λ_2, that describes the deviation from expected Poisson dispersion, or the ratio of the variance of the distribution to the mean.

> NB: A Poisson distribution has a variance-to-mean ratio equal to 1. As an exercise, show that the variance of the spider count distribution is approximately equal to its mean.

$\lambda_2 > 0$ implies that organisms are over-dispersed and may be attracted to one another, forming more large-number clusters than expected by chance. $\lambda_2 < 0$ implies that organisms are under-dispersed and are likely repulsed by one another, resulting in a more even distribution across space and, hence, lower count variance. The LGP distribution, in terms of these two paramters, is

$P(k) = \lambda_1(\lambda_1 + k\lambda_2)^{k-1}\frac{e^{-(\lambda_1+k\lambda_2)}}{k!}$.

When $\lambda_2 = 0$, we get precisely the Poisson distribution. Back to the sowbugs! From an empiral fit, we find that $\lambda_2 = 0.53214$ produces the best agreement between this model and the data. A simple estimate of λ_1 follows from

$\lambda_1 = \text{mean}(1 - \lambda_2)$.

```
mean_sowbugs <- sum(x$k_number_of_arthropods*x$C_count_of_boards_with_k_sowbugs) / sum(x$C_count_of_boa:
lambda2_sowbugs = 0.53214
lambda1_sowbugs = mean_sowbugs*(1-lambda2_sowbugs)
```

We can now use these parameters to show how well this model fits the sowbug data. First let's load a package that contains a function for the LGP equation, that returns a count distribution over a range of k values, given input parameters λ_1 and λ_2.

```
library(RMKdiscrete)
LGP_k_sowbugs <- data.frame(k=k, probability=dLGP(k,lambda1_sowbugs,lambda2_sowbugs))
```

> **Exercise**
>
> Add a line to your sowbug plot, showing the output of the LGP function with our defined parameters. How well does this model fit the data?

Weevils are antisocial

In 1975, Rodger Mitchell reported on data from an experiment on weevil egg-laying behavior. Weevils lay their eggs in mung beans, then the larva hatches within the bean and eats the bean as it matures. A larger bean with fewer other hatching eggs to compete with will increase the larva's chance of survival. Mitchell measured how many eggs were laid per bean in an experiment were eggs were plentiful and competition was low. Load the bean weevil data

```
weevil_data <- read.delim("../data/mitchell_weevil_egg_data_1975.csv", sep=',')
N_eggs <- sum(weevil_data$C_count_of_beans_with_k_eggs)
weevil_data$P_k_weevils <- weevil_data$C_count_of_beans_with_k_eggs/N_eggs
```

...and plot it.

```
g_weevils <- ggplot(weevil_data, aes(x=k_number_of_eggs, y=P_k_weevils)) + geom_point(size = 3) + xlab(
g_weevils
```

Exercise

Add a line to your plot that shows a Poisson distribution with the same mean number of eggs per bean. How does the Poisson fit compare to the data? Is the variance higher or lower, by eye, than for the Poisson distribution?

With these data, it's possible to estimate both λ_1 and λ_2 from the data and get a good fit of the LGP. You can estimate

$\lambda_2 = 1 - (\frac{\text{variance}}{\text{mean}})^{-1/2}$, and $\lambda_1 = \text{mean}(1 - \lambda_2)$.

Exercise

Compute these two parameters. Use `dLPG` and add a line to your weevil plot for these results. Does this distribution fit the data well?

Introduction to version control

Version control is a way to record and organize changes to a set of files and directories (e.g., the directory containing one of your projects). Over time, a version control system (VCS) builds a database storing all the changes you perform (a *repository*), making the whole history of the project available.

When you start working on a new project, you tell your VCS to keep track of all the changes, additions, and deletions. At any point, you can *commit* the changes, effectively building a snapshot of the project in the repository. This snapshot of the project is then accessible—you can recover previously committed versions of files, including metadata such as who changed what, when and why. Also, you can easily start tracking files for an existing project.

Version control is especially important for collaborative projects: everybody can simultaneously work on the project, even on the same file. Conflicting changes are reported, and can be managed using side-by-side comparisons. The possibility of *branching* allows you to experiment with changes (e.g., shall we rewrite the introduction of this paper?), and then decide whether to *merge* them into the project.

Why use Version Control?

Version control is fundamental to keeping a project tidily organized and a central piece to making scientific computing as automated, reproducible, and easy to read and share as possible.

Many scientists keep backup versions of the same project over time or append a date/initials to different versions of the same file (e.g., various drafts of the same manuscript). This manual approach quickly becomes unmanageable, with simply too many files and versions to keep track of in a timely and organized manner. Version control allows you to access all previously committed versions of the files and directories of your project. This means that it is quite easy to undo short-term changes: Bad day? Just go back to yesterday's version! You can also access previous stages of the project: "I need to access the manuscript's version and all the analysis files in exactly the state that they where in when I sent the draft for review three months ago." Checking out an entire project at a certain point in time is easy with a version control system but much more difficult with `Dropbox` or `Google Drive`.

Version control is useful for small, and essential for large collaborative projects. It vastly improves the workflow, efficiency and reproducibility. Without it, it is quite easy to lose track of the status of a manuscript

(who has the most recent version?), or lose time (I cannot proceed with my changes, because I have to wait for the edits of my collaborators).

Version control might look like overkill at first. However, with a little bit of practice you will automatically run a few commands before you start working on a project, and again once you are done working on it. A small price to pay, considering the advantages. Simply put, using version control makes you a more organized and efficient scientist.

> Stefano's testimonial: "Our laboratory adopted version control for all our projects in 2010, and sometimes we wonder how we managed without it."

Throughout this introduction, we illustrate the advantages of using version control for conducting scientific research, and assume that each repository is a scientific project. At first, we work with *local* repositories, meaning that all the files are stored exclusively on your computer. Then, we introduce *remote* repositories, which are also hosted on a web server, making it easy for you to share your projects with others (or work on the same project from different computers).

Git

For this introduction to version control, we use `git`, which is one of the most popular version control systems. `git` is also free software, and is available for all computer architectures. Many good tutorials are available online, and many websites will host your repositories for free.

Other options you might want to consider are `Mercurial` (very similar to `git`) and Subversion Version Control (`svn`), which is an older system, but still widespread.

There are two main paradigms for VCSs allowing multiple users to collaborate: in a *centralized* VCS (e.g., `svn`), the whole history of a project is stored exclusively on a server, and users download the most current snapshot of the project, modify it, and send the changes to the server; in a *distributed* VCS (e.g., `git`), the complete history of the repository is stored on each user's computer.

`git` was initially developed by Linus Torvalds (the "Linu" in Linux), exactly for the development of the Linux kernel. It was first released in 2005 and has since become the most widely adopted version control system.

Configuring git

First, open a terminal window.

The first time you use `git`, (or whenever you install `git` on a new computer), you need to set up the environment. To store your preferred user name and email, type:

```
$ git config --global user.name "Charles Darwin"
$ git config --global user.email crdarwin@royalsociety.org
```

Optionally, you can set up your preferred text editor (e.g., `gedit` or `emacs`), which will be used to write the messages associated wih your commits:

```
$ git config --global core.editor gedit
```

To check all of your settings and see all available options, type:

```
$ git config --list
```

How to get help in Git

For a brief overview of common `git` commands, open your terminal and type:

```
$ git help
```

Now try typing:

```
$ man git
```

you will see that the name of the program is `git -  the stupid content tracker`. Git behaves stupidly in the sense that the system tracks any content without being selective (which can be a good and bad thing). The manual page contains a description of all the commands but it is much more pleasing to read them online.

Everyday Git

To illustrate the basic operations in `git`, we consider the case of a local repository: all the versions are stored only in your computer, and we assume that you are the only person working on the project. Once familiarized with the basics of `git`, we introduce the use of remote repositories for collaborative projects.

Workflow

Here is the typical day in the life of a git user:

> When starting a new project, create a directory using terminal and initialize a repository. Start working on your files until you reach a milestone or until you are done for the day. Check what has changed since your last snapshot. Decide which files to add to the new snapshot of the project (selected files, or everything). Create the snapshot by committing your changes, including a detailed description. Start changing the files again, add them to the new snapshot, and commit.

As you can see, there are only a few commands that you need to master for everyday work. Let's try our hands at this workflow and create a simple repository. Open a terminal window and create a new directory called `git_test`.

```
$ cd /git_test
$ mkdir originspecies
$ cd originspecies
$ git init
```

where we have moved to a newly created directory "originspecies", and run the command `git init`, which initializes an empty repository. You can set up a `git` repository for an existing project by changing to the directory containing the project and typing `git init` in the terminal. We recommend to always run `pwd` before initializing a `git` repository to confirm that you are in the correct directory. The last thing you want is to track changes to your entire computer because you happened to be in your root directory.

We can check the status of the repository by running:

```
$ git status
```

Now we create our first file:

```
$ touch origin.txt
```

and then start editing it. We could use a text editor, but for the moment let's stick to the command:

```
$ echo "An abstract of an Essay on the Origin of Species and Varieties Through natural selection" > ori
```

and check that our command went through:

```
$ cat origin.txt
```

`git` does not track any file unless you tell it to do so. We can set the file `origin.txt` for tracking with the command:

```
$ git add origin.txt
```

We can see that the status of the repository has changed:

```
$ git status
```

Every time we want to signal `git` that a file needs to be tracked, we need to add it to the index of files in the repository, using the command `git add FILENAME`. If we want to simply add all the files in a directory and subdirectories, use `git add .` (where the "." means "current directory"). To add all the files contained in the repository directory (including "hidden" files), use `git add --all`.

Note that `git` and similar systems are ideally suited to work with text files: `.csv` and `.txt` files, code (`.R`, `.py`, etc.), LaTeX manuscripts (`.tex`, `.bib`), etc. If you track *binary* files, `git` will simply save a new version any time you change it, but you will not be able to automatically see the differences between different versions of the files.

Once we are finished creating, deleting and modifying our files, we can create a snapshot of the project by *committing* the changes. When should one commit? `git`'s motto is "commit early, commit often". Every time you commit your changes, these are permanently saved in the repository. Ideally, every commit should represent a meaningful step on the path to completing the project (examples: "drafted introduction", "implemented simulation", "rewritten hill-climber", "added the references", ...). As a rule of thumb, you should commit every time you can explain the meaning of your changes in a few lines.

Now it's time to perform our first commit:

```
$ git commit -m "started the book"
```

where `-m` stands for "message". If no message is entered, `git` will open your default text editor, so that you can write a longer message detailing your changes. These messages are very important! In fact, you can use them to navigate and understand the history of your project. Make sure you always spend a few extra seconds to document what you did by writing a meaningful and detailed message. In a few months, you will have forgotten all the details associated with a commit, so that your future self will be grateful for detailed messages specifying how and why your project has changed.

The history of the repository can be accessed by typing:

```
$ git log
```

The long line of numbers and letters after the word `commit` is the "checksum" associated with the commit (i.e., a number `git` uses to make sure that all changes have been successfully stored). You can think of this number as the "fingerprint" of the commit.

Now, let's change a tracked file:

```
echo "On the Origin of Species, by Means of Natural Selection, or the Preservation of Favoured Races in
```

a much more powerful, albeit a bit long, title. The repository has changed and we can investigate the changes using `status`:

```
$ git status
```

showing that a) a file that is being tracked has changed, and b) that these changes have not been *staged* (i.e., marked) to be committed yet. You can keep modifying the file, and once you are satisfied with it, use *git add* to make these changes part of the next commit. For example:

```
$ git add .
$ git commit -m "Changed the title as suggested by Murray"
```

You now see the new commit in the history of the repository:

```
$ git log
```

That's it! For 99% of your time working with `git`, all you need to do is to follow these simple steps:

- When you are creating a new project

```
$ mkdir newproject
$ cd newproject
$ git init # initialize repository
```

- Daily routine:
- 1) Change the files
- 2) Check status

```
$ git status
```

- 3) Add files to the snapshot

```
$ git add --all
```

- 4) Commit the snapshot

```
$ git commit -m "my descriptive message"
```

Besides these basic commands, `git` allows you to do much more. We explore a few of the more advanced features in the next sections.

Moving and removing files

When you want to move or remove files or directories under version control, you should let `git` know, so that it can update the index of files to be tracked. Doing so is quite easy: simply put `git` in front of the command you would run to perform the operation:

```
$ git rm filetorem.txt
$ git rm *.txt
$ git mv myoldname.txt mynewname.csv
```

Deleting changes and reverting to last commit

Sometimes, changes can go horribly wrong, such that you would like to scrap everything, and abandon all the changes you made to one (or more) file(s). This can be accomplished by typing:

```
$ git checkout filetoreset.txt
```

once you run this command, the file will be in the last version you committed. Use this command with caution, as the changes you made are irrevocably lost.

Remote repositories

So far, we have been working with local repositories, hosted in only one computer. When collaborating (or when working on the same project from different computers), it is convenient to host the repository on a server, so that the collaborators can see your changes, and add their own.

There are two main options: a) setup your own server (this is beyond the scope of this introduction, but see Section **??** for pointers on how to do this); b) host your repositories on a website offering this service (for a fee, or for free).

The most popular option for hosting Open Source projects, where the whole world can see what and when you commit, is GitHub, which will store all your public repositories for free.

After the initial setup (which is specific to the service you use, and thus not treated here), you only need to add two new commands to your `git` workflow: `pull` and `push`. When you want to work on a project that is tracked by a remote repository, you `pull` the most recent version from the server and work on your local copy using the commands illustrated above. When you are done, you `push` your commits to the server so that other users can see them. Your workflow will look like:

- Get the latest version of the project

```
$ git pull
```

- Start making changes. When done run:

```
$ git add --all
```

- Time to commit:

```
$ git commit -m "A meaningful message"
```

- Rinse and repeat...
- Ok, everyone should see my changes

```
$ git push
```

Note that the interaction with the server only happens when you `pull` and `push`. In between these commands, you are in control, and can decide when and what to communicate to your collaborators.

Branching and merging

Most VCSs allow you to *branch* from the main development of the project (i.e., experiment freely with your project without messing up the original version). It is like saying "Save as" with a different name, but with the possibility of easily merging the two parallel versions of the file. Typical examples of a branching point in scientific projects are: a) you want to try a different angle for the Introduction of your manuscript, but you are not sure it's going to be better; b) you want to try to rewrite a piece of code, to see whether it will be faster, but surely you do not want to ruin the current version that is working just fine; c) you want to keep working on the figures, while your collaborator is editing the manuscript. In these cases, you are working on an "experimental feature", while leaving the main project unaltered (or while other people are working on it). Once you are satisfied with your changes, you would like to *merge* them with the main version of the project.

In `git`, creating a new branch is quite easy. Suppose that you are working on a project, and that you have never branched before. Suppose that you want to try amending a piece of the code (for example, to make it run faster), but don't want to touch the code that is already working. You need to *branch*:

```
$ git branch fastercode
```

creates the new branch `fastercode`.

- List branches in your repository

```
$ git branch
```

- Switch branch

```
$ git checkout fastercode
```

which means that now you can experiment freely. Suppose that you edit the code, and commit the changes:

- Change the code, and then commit

```
$ git commit -m "Managed to make code faster"
```

To summarize, here's how you work with branches:

- Create a new branch

```
$ git branch mybranch
```

- See that you are in master

```
$ git branch
```

- Move to new branch

```
$ git checkout mybranch
```

- Start working on the branch
- . . .
- Add and commit changes

```
$ git add --all
$ git commit -m "my message"
```

- Once you are satisfied, and want to include the changes into the main project, move back to main trunk

```
$ git checkout master
```

- Merge the branch

```
$ git merge mybranch -m "message for merge"
```

- [optional] delete the branch

```
$ git -d mybranch
```

References and readings

Journal articles

All of the data used in this tutorial come from original research papers that are in the `readings` folder. Also in the `readings` folder, you will find an article by Roger Peng arguing for setting standards for reproducible research in computational science. It's a short article that we hope you will read and adopt as best practices for your own work.

Books and tutorials

There are very many good books and tutorials on `git`. We are particularly fond of *Pro Git*, by Scott Chacon and Ben Straub. You can either buy a physical copy of the book, or read it online for free.

Both GitHub and Atlassian (managing Bitbucket) have their own tutorials.

A great way to try out Git in 15 minutes is here.

Software Carpentry offers intensive on-site workshops and online tutorials.

GUIs

There are several Graphical User Interfaces for git.

Data Challenge

This tutorial has been focused on making your code reproducible and sharable with the wider world. Many journal now *require* that you submit your code and data when you publish your paper. One such article was recently written on the False Discovery Rate in medical research (Jaeger and Leek (2014) can be found in the `readings` folder). While an earlier article claimed that over 80% of medical results were probably false discoveries, this work uses a slightly more sophisticated analysis of the p-values reported in the medical literature, to determine whether the observed distribution of p-values is significantly different from those drawn from a uniform distribution. Reassuringly, the authors find that the false discovery rate is lower (~20%), but still troubling, across many medical journals.

- 1) Use the published code in: Jager, Leah R., and Jeffrey T. Leek. "An estimate of the science-wise false discovery rate and application to the top medical literature." Biostatistics 15.1 (2013): 1-12. . . . to reproduce the results of the paper. In particular, report the mean and standard deviaiont of the FDR of the journal The Lancet to 4 significant digits.

- 2) Use the scraped abstract data provided in `data` and use the authors' code to compute the FDR for the journal PNAS. Find the mean and s.d. of the FDR for this journal.

- 3) Bonus question: use the authors' code to first scrape PubMed for abstracts, and next compute the FDR for the Journal of Psychiatric Neuroscience. You will find a bug in the code that you will need to hunt down and fix.

- 4) Bonus question: Find the part of the code where the authors' initialize their random number generator.

Go to URL to submit your team's answers.